

ISA 401/501: Business Intelligence & Data Visualization

25: A Short Introduction to Clustering

Fadel M. Megahed, PhD

Enders Associate Professor
Farmer School of Business
Miami University

 @FadelMegahed


 fmegahed

 fmegahed@miamioh.edu


 Automated Scheduler for Office Hours

Fall 2022


A Recap of What we Learned Last Class

- Describe the goals & functions of data mining
- Understand the statistical limits on data mining
- Describe the data mining process
- What is “frequent itemsets” & the application of this concept
- Explain how and why “association rules” are constructed
- Use  to populate both concepts

Kahoot: A Recap of Phase 3 of Class So Far

Let us go to Kahoot and compete for a \$10  Starbucks gift card. To evaluate your understanding of the material, please answer the questions correctly and as quickly as possible to get the most points.

Learning Objectives for Today's Class

- Describe the different steps of the k-means algorithm
- Cluster using k-means (by hand)
- Cluster using k-means (software)
 - 
 - Tableau


An Overview of Clustering Techniques

The Problem of Clustering


- Given a **set of (high-dimensional) observations**, with a notion of **distance** between observations, **group the observations** into **some number of clusters**, so that:
 - Members of a cluster are close/similar to each other
 - Members of different clusters are dissimilar
- **Usually:**
 - The observations are in a high-dimensional space
 - Similarity is defined using a distance measure, e.g.,
 - Euclidean, Cosine, Jaccard, edit distance, etc.

Clustering in 2D Space

Meet the Palmer penguins



Anatomical description of the dataset:




Clustering in 2D Space: Formulation

- Given a **set of observations (each containing bill length and depth)**, with a notion of **Euclidean distance** between observations, **group the observations** into **3 clusters**, so that:
 - Members of a cluster are close/similar to each other
 - Members of different clusters are dissimilar
- Note we are assuming that we did not have a "label/type" for each penguin.

Clustering in 2D Space: Raw Data

Penguin bill dimensions


Bill length and depth for Adelie, Chinstrap and Gentoo Penguins at Palmer Station LTER



Clustering in 2D Space: Labeled Raw Data

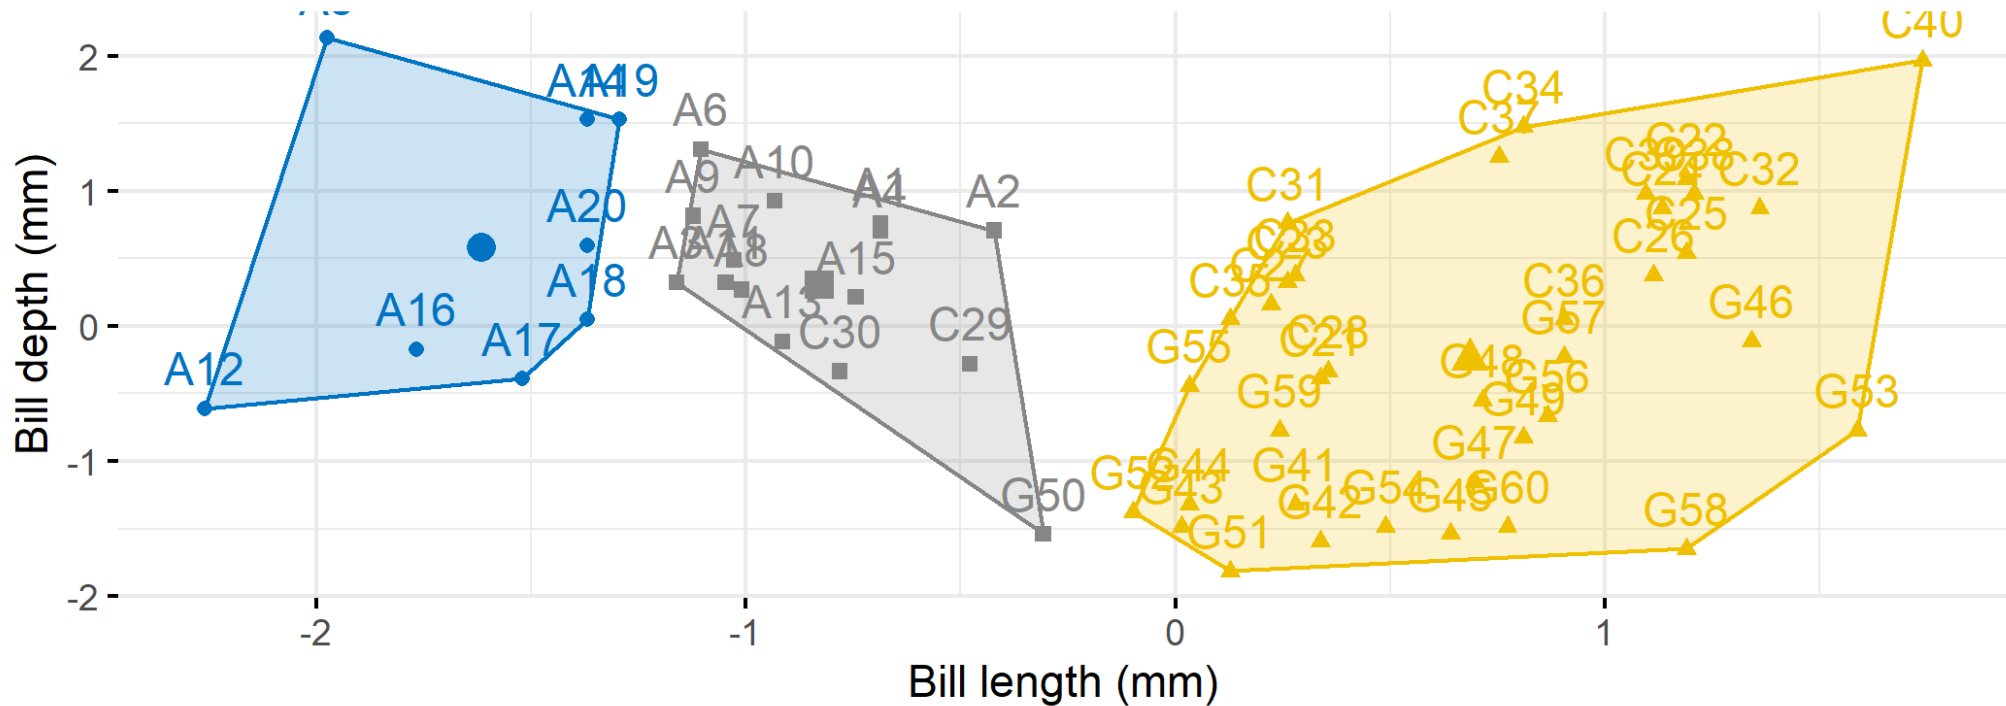
Penguin bill dimensions

Bill length and depth for Adelie, Chinstrap and Gentoo Penguins at Palmer Station LTER



Clustering in 2D Space: Clustering Results

Clustering of a sample of 60 penguins into three groups
Using only bill length and depth



Comments on the 2D Clustering Problem

Even though the 2D Space clustering problem is the easiest problem to "solve" since we can benefit by plotting the data, **clustering is hard**.

Some important questions:

- With all the variables being numerical, we often assume **Euclidean distance**. This can be problematic when:
 - variables have significantly different scales
 - we are including information that is not pertinent to grouping
- How do you determine the number of clusters (k)?
- How to represent a cluster of many points?
- How do we determine the "nearness" of clusters?

An Overview of Clustering Methods

Categories	Abb. name	Volume			Variety		Velocity	Other criterion
		Size of Dataset	Handling High Dimensionality	Handling Noisy Data	Type of Dataset	Clusters Shape	complexity of Algorithm	Input Parameter
Partitional algorithms	K-Means [25]	Large	No	No	Numerical	Non-convex	$O(nkd)$	1
	K-modes [19]	Large	Yes	No	Categorical	Non-convex	$O(n)$	1
	K-medoids [33]	Small	Yes	Yes	Categorical	Non-convex	$O(n^2dt)$	1
	PAM [31]	Small	No	No	Numerical	Non-convex	$O(k(n-k)^2)$	1
	CLARA [23]	Large	No	No	Numerical	Non-convex	$O(k(40+k)^2+k(n-k))$	1
	CLARANS [32]	Large	No	No	Numerical	Non-convex	$O(kn^2)$	2
	FCM [6]	Large	No	No	Numerical	Non-convex	$O(n)$	1
Hierarchical algorithms	BIRCH [40]	Large	No	No	Numerical	Non-convex	$O(n)$	2
	CURE [14]	Large	Yes	Yes	Numerical	Arbitrary	$O(n^2 \log n)$	2
	ROCK [15]	Large	No	No	Categorical and Numerical	Arbitrary	$O(n^2 + nmmma + n^2 \log n)$	1
	Chameleon [22]	Large	Yes	No	All type of data	Arbitrary	$O(n^2)$	3
	ECHIDNA [26]	Large	No	No	Multivariate Data	Non-convex	$O(N * B(1 + \log_B m))$	2
Density-based algorithms	DBSCAN [9]	Large	No	No	Numerical	Arbitrary	$O(n \log n)$ If a spatial index is used Otherwise, it is $O(n^2)$.	2
	OPTICS [5]	Large	No	Yes	Numerical	Arbitrary	$O(n \log n)$	2
	DBCLASD [39]	Large	No	Yes	Numerical	Arbitrary	$O(3n^2)$	No
	DENCLUE [17]	Large	Yes	Yes	Numerical	Arbitrary	$O(\log D)$	2
Grid- based algorithms	Wave-Cluster [34]	Large	No	Yes	Special data	Arbitrary	$O(n)$	3
	STING [37]	Large	No	Yes	Special data	Arbitrary	$O(k)$	1
	CLIQUE [21]	Large	Yes	No	Numerical	Arbitrary	$O(Ck + mk)$	2
	OptiGrid [18]	Large	Yes	Yes	Special data	Arbitrary	Between $O(nd)$ and $O(nd \log n)$	3
Model- based algorithms	EM [8]	Large	Yes	No	Special data	Non-convex	$O(knp)$	3
	COBWEB [12]	Small	No	No	Numerical	Non-convex	$O(n^2)$	1
	CLASSIT [13]	Small	No	No	Numerical	Non-convex	$O(n^2)$	1
	SOMs [24]	Small	Yes	No	Multivariate Data	Non-convex	$O(n^2m)$	2

k-means Algorithm

General Idea

The K-Means algorithm clusters data by trying to separate samples in n groups of equal variance, minimizing a criterion known as the **inertia** or **within-cluster sum-of-squares** (see below). This algorithm requires the **number of clusters to be specified**.

$$\sum_{i=0}^n \min_{\mu_j \in C} (||x_i - \mu_j||^2)$$

Inertia is a measure of how internally coherent clusters are; however, it suffers from various drawbacks:

- Inertia makes the assumption that clusters are convex and isotropic, which is not always the case. It responds poorly to elongated clusters, or manifolds with irregular shapes.
- Inertia is not a normalized metric: we just know that lower values are better and zero is optimal. But in very high-dimensional spaces, Euclidean distances tend to become inflated.

The Steps of the K-Means Algorithm

In basic terms, the algorithm has three steps.

- Step 0 chooses the initial centroids, with the most basic method being to choose k samples from the dataset X . After initialization, K-means consists of looping between the remaining two steps.
- Step 1 assigns each sample to its nearest centroid.
- Step 2 creates new centroids by taking the mean value of all of the samples assigned to each previous centroid. The difference between the old and the new centroids are computed.

The algorithm repeats these last two steps the centroids do not move significantly.

Out-of-Class Activity

Use the k-means algorithm to cluster the following observations. Use $k = 2$ and Euclidean distance. **Use [this handout](#) to go through the k-means algorithm implementation (by hand).**

Observation	X1	X2
1	1.0	1.0
2	1.5	2.0
3	3.0	4.0
4	5.0	7.0
5	3.5	5.0
6	4.5	5.0
7	3.5	4.5

Practical Issues with K-Means Clustering

Data	Prep	K-means (k=3)	Optimal k	Viz Clusters
<pre>pacman::p_load(tidyverse, palmerpenguins, magrittr) penguins_tbl = penguins # our data for today penguins_tbl # printing it out</pre>				
<pre>## # A tibble: 344 × 8 ## species island bill_length_mm bill_depth_mm flipper_...¹ body_...² sex year ## <fct> <fct> <dbl> <dbl> <int> <int> <fct> <int> ## 1 Adelie Torgersen 39.1 18.7 181 3750 male 2007 ## 2 Adelie Torgersen 39.5 17.4 186 3800 fema... 2007 ## 3 Adelie Torgersen 40.3 18 195 3250 fema... 2007 ## 4 Adelie Torgersen NA NA NA NA <NA> 2007 ## 5 Adelie Torgersen 36.7 19.3 193 3450 fema... 2007 ## 6 Adelie Torgersen 39.3 20.6 190 3650 male 2007 ## 7 Adelie Torgersen 38.9 17.8 181 3625 fema... 2007 ## 8 Adelie Torgersen 39.2 19.6 195 4675 male 2007 ## 9 Adelie Torgersen 34.1 18.1 193 3475 <NA> 2007 ## 10 Adelie Torgersen 42 20.2 190 4250 <NA> 2007 ## # ... with 334 more rows, and abbreviated variable names ¹flipper_length_mm, ## # ²body_mass_g</pre>				

Practical Issues with K-Means Clustering

Data	Prep	K-means (k=3)	Optimal k	Viz Clusters
------	------	---------------	-----------	--------------

```
penguins_tbl %<>%  
  # selecting relevant cols  
  select(species, bill_length_mm, bill_depth_mm, flipper_length_mm, body_mass_g) %>%  
  na.omit() %>% # removing NAs  
  mutate_at(vars(-species), scale) # scaling numeric variables  
  
penguins_tbl # printing it out
```

```
## # A tibble: 342 × 5  
##   species bill_length_mm[,1] bill_depth_mm[,1] flipper_length_mm[,1] body_mas...1  
##   <fct>          <dbl>          <dbl>          <dbl>          <dbl>  
## 1 Adelie        -0.883            0.784          -1.42         -0.563  
## 2 Adelie        -0.810            0.126          -1.06         -0.501  
## 3 Adelie        -0.663            0.430          -0.421        -1.19  
## 4 Adelie        -1.32             1.09          -0.563        -0.937  
## 5 Adelie        -0.847            1.75          -0.776        -0.688  
## 6 Adelie        -0.920            0.329          -1.42         -0.719  
## 7 Adelie        -0.865            1.24          -0.421         0.590  
## 8 Adelie        -1.80             0.480          -0.563        -0.906  
## 9 Adelie        -0.352            1.54          -0.776         0.0602  
## 10 Adelie       -1.12           -0.0259          -1.06         -1.12  
## # ... with 332 more rows, and abbreviated variable name 1body_mass_g[,1]
```

Practical Issues with K-Means Clustering

Data	Prep	K-means (k=3)	Optimal k	Viz Clusters
------	------	---------------	-----------	--------------

```
km_res = kmeans(x = penguins_tbl %>% select(-species), # input data with no label
                centers = 3) # k =3

# tabulating the results with rows corresponding to true labels and the columns corresponding to clusters
table(penguins_tbl$species, km_res$cluster)
```

```
##
##      1    2    3
## Adelie  0    0 151
## Chinstrap 0    1  67
## Gentoo 66   57   0
```

Practical Issues with K-Means Clustering

Data

Prep

K-means (k=3)

Optimal k

Viz Clusters

```
pacman::p_load(NbClust)

km_res_nbclust = NbClust(
  data = penguins_tbl %>% select(-species),
  distance = "euclidean",
  min.nc = 2, max.nc = 10,
  method = "kmeans", index = "all")

table(penguins_tbl$species, km_res_nbclust$Best.partition)
```

```
## *** : The Hubert index is a graphical method of determining the nu
##           In the plot of Hubert index, we seek a significant
##           significant increase of the value of the measure i
##           index second differences plot.
```

```
## *** : The D index is a graphical method of determining the number
##           In the plot of D index, we seek a significant knee
##           second differences plot) that corresponds to a sig
##           the measure.
## *****
## * Among all indices:
## * 8 proposed 2 as the best number of clusters
## * 11 proposed 3 as the best number of clusters
## * 1 proposed 4 as the best number of clusters
## * 3 proposed 5 as the best number of clusters
## * 1 proposed 10 as the best number of clusters
## ***** Conclusion *****
## * According to the majority rule, the best number of clusters is
```

Practical Issues with K-Means Clustering

Data

Prep


K-means (k=3)

Optimal k

Viz Clusters

```
pacman::p_load(factoextra)

fviz_cluster(object =
  list(
    cluster = km_res_nbclust$Best.partition,
    data = penguins_tbl %>% select(-species)
  ),
  ellipse.type = "convex",
  palette = "jco",
  ggtheme = theme_minimal()
)
```



Summary of Practical Issues


- Rescale numeric data prior to k-means implementation. The scaling can be:
 - a z-transformation similar to what we did in the example
 - a 0-1 scaling
 - converting count data into percentage or counts per a certain number of the population
 - etc.
- Use more than one metric to determine k when using k-means clustering
- Your cluster solution is not the end result, you will need to:
 - visualize it in appropriate way (simple representation as in the previous slide, [spatially](#), [time-based](#), etc.)
 - Attempt to explain the cluster membership using an appropriate binomial/multinomial model (e.g., see [this analysis](#))

k-means in Tableau

Let us use Tableau to implement the k-means clustering implementation on the 60 sample observations from the penguins dataset as shown in Slide 7 of this presentation.

Recap

Summary of Main Points

- Describe the different steps of the k-means algorithm
- Cluster using k-means (by hand)
- Cluster using k-means (software)
 - 
 - Tableau